

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler #for feature scaling
from sklearn.model_selection import train_test_split #train-test

#for simple linear regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import PolynomialFeatures

```

data exprolartion

```

df =
pd.read_csv('/kaggle/input/auto-scout-car-price/final_scout_not_dummy.
csv')
print("Shape of the dataset:", df.shape)

```

Shape of the dataset: (15915, 23)

df.head()

	make_model	body_type	price		vat	km	Type
Fuel \							
0	Audi A1	Sedans	15770	VAT deductible	56013.0	Used	
Diesel							
1	Audi A1	Sedans	14500	Price negotiable	80000.0	Used	
Benzine							
2	Audi A1	Sedans	14640	VAT deductible	83450.0	Used	
Diesel							
3	Audi A1	Sedans	14500	VAT deductible	73000.0	Used	
Diesel							
4	Audi A1	Sedans	16790	VAT deductible	16200.0	Used	
Diesel							

	Gears		Comfort_Convenience	\
0	7.0	Air conditioning,Armrest,Automatic climate con...		
1	7.0	Air conditioning,Automatic climate control,Hil...		
2	7.0	Air conditioning,Cruise control,Electrical sid...		
3	6.0	Air suspension,Armrest,Auxiliary heating,Elect...		
4	7.0	Air conditioning,Armrest,Automatic climate con...		

		Entertainment_Media	...
Previous_Owners	\		
0	Bluetooth,Hands-free equipment,On-board comput...		...
2.0			

```

1 Bluetooth,Hands-free equipment,On-board comput... ...
1.0
2 MP3,On-board computer ...
1.0
3 Bluetooth,CD player,Hands-free equipment,MP3,0... ...
1.0
4 Bluetooth,CD player,Hands-free equipment,MP3,0... ...
1.0

```

	hp_kw	Inspection_new	Paint_Type	Upholstery_type	Gearing_Type	\
0	66.0	1	Metallic	Cloth	Automatic	
1	141.0	0	Metallic	Cloth	Automatic	
2	85.0	0	Metallic	Cloth	Automatic	
3	66.0	0	Metallic	Cloth	Automatic	
4	66.0	1	Metallic	Cloth	Automatic	

	Displacement_cc	Weight_kg	Drive_chain	cons_comb
0	1422.0	1220.0	front	3.8
1	1798.0	1255.0	front	5.6
2	1598.0	1135.0	front	3.8
3	1422.0	1195.0	front	3.8
4	1422.0	1135.0	front	4.1

```
[5 rows x 23 columns]
```

```

df.info()
df.describe()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15915 entries, 0 to 15914
Data columns (total 23 columns):

```

#	Column	Non-Null Count	Dtype
0	make_model	15915 non-null	object
1	body_type	15915 non-null	object
2	price	15915 non-null	int64
3	vat	15915 non-null	object
4	km	15915 non-null	float64
5	Type	15915 non-null	object
6	Fuel	15915 non-null	object
7	Gears	15915 non-null	float64
8	Comfort_Convenience	15915 non-null	object
9	Entertainment_Media	15915 non-null	object
10	Extras	15915 non-null	object
11	Safety_Security	15915 non-null	object
12	age	15915 non-null	float64
13	Previous_Owners	15915 non-null	float64
14	hp_kw	15915 non-null	float64
15	Inspection_new	15915 non-null	int64
16	Paint_Type	15915 non-null	object

```

17 Upholstery_type      15915 non-null object
18 Gearing_Type         15915 non-null object
19 Displacement_cc      15915 non-null float64
20 Weight_kg            15915 non-null float64
21 Drive_chain          15915 non-null object
22 cons_comb            15915 non-null float64
dtypes: float64(8), int64(2), object(13)
memory usage: 2.8+ MB

```

	price	km	Gears	age \
count	15915.000000	15915.000000	15915.000000	15915.000000
mean	18024.380584	32089.995708	5.937355	1.389695
std	7381.679318	36977.214964	0.704772	1.121306
min	4950.000000	0.000000	5.000000	0.000000
25%	12850.000000	1920.500000	5.000000	0.000000
50%	16900.000000	20413.000000	6.000000	1.000000
75%	21900.000000	46900.000000	6.000000	2.000000
max	74600.000000	317000.000000	8.000000	3.000000

	Previous_Owners	hp_kW	Inspection_new	Displacement_cc
\count	15915.000000	15915.000000	15915.000000	15915.000000
mean	1.042853	88.499340	0.247063	1428.661891
std	0.339178	26.674341	0.431317	275.804272
min	0.000000	40.000000	0.000000	890.000000
25%	1.000000	66.000000	0.000000	1229.000000
50%	1.000000	85.000000	0.000000	1461.000000
75%	1.000000	103.000000	0.000000	1598.000000
max	4.000000	294.000000	1.000000	2967.000000

	Weight_kg	cons_comb
count	15915.000000	15915.000000
mean	1337.700534	4.832124
std	199.682385	0.867530
min	840.000000	3.000000
25%	1165.000000	4.100000
50%	1295.000000	4.800000
75%	1472.000000	5.400000
max	2471.000000	9.100000

data cleaning

```
df.isnull().sum()
```

make_model	0
body_type	0
price	0
vat	0
km	0
Type	0
Fuel	0
Gears	0
Comfort_Convenience	0
Entertainment_Media	0
Extras	0
Safety_Security	0
age	0
Previous_Owners	0
hp_kW	0
Inspection_new	0
Paint_Type	0
Upholstery_type	0
Gearing_Type	0
Displacement_cc	0
Weight_kg	0
Drive_chain	0
cons_comb	0

dtype: int64

```
df.isna().sum()
```

make_model	0
body_type	0
price	0
vat	0
km	0
Type	0
Fuel	0
Gears	0
Comfort_Convenience	0
Entertainment_Media	0
Extras	0
Safety_Security	0
age	0
Previous_Owners	0
hp_kW	0
Inspection_new	0
Paint_Type	0
Upholstery_type	0
Gearing_Type	0

```

Displacement_cc      0
Weight_kg             0
Drive_chain           0
cons_comb             0
dtype: int64

df.duplicated().sum()

1673

cleaned_df=df.drop_duplicates()
cleaned_df.duplicated().sum()

0

#need only the numerical values
cleaned_df=df.select_dtypes(include=[np.number])
cleaned_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15915 entries, 0 to 15914
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   price                 15915 non-null  int64   
1   km                   15915 non-null  float64  
2   Gears                 15915 non-null  float64  
3   age                  15915 non-null  float64  
4   Previous_Owners      15915 non-null  float64  
5   hp_kW                15915 non-null  float64  
6   Inspection_new       15915 non-null  int64   
7   Displacement_cc      15915 non-null  float64  
8   Weight_kg            15915 non-null  float64  
9   cons_comb            15915 non-null  float64  
dtypes: float64(8), int64(2)
memory usage: 1.2 MB

```

data visualization

```

# Now create the pairplot
sns.pairplot(cleaned_df, kind='scatter', plot_kws={'s': 50, 'alpha': 0.5})

/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
  with pd.option_context('mode.use_inf_as_na', True):
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:

```

FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

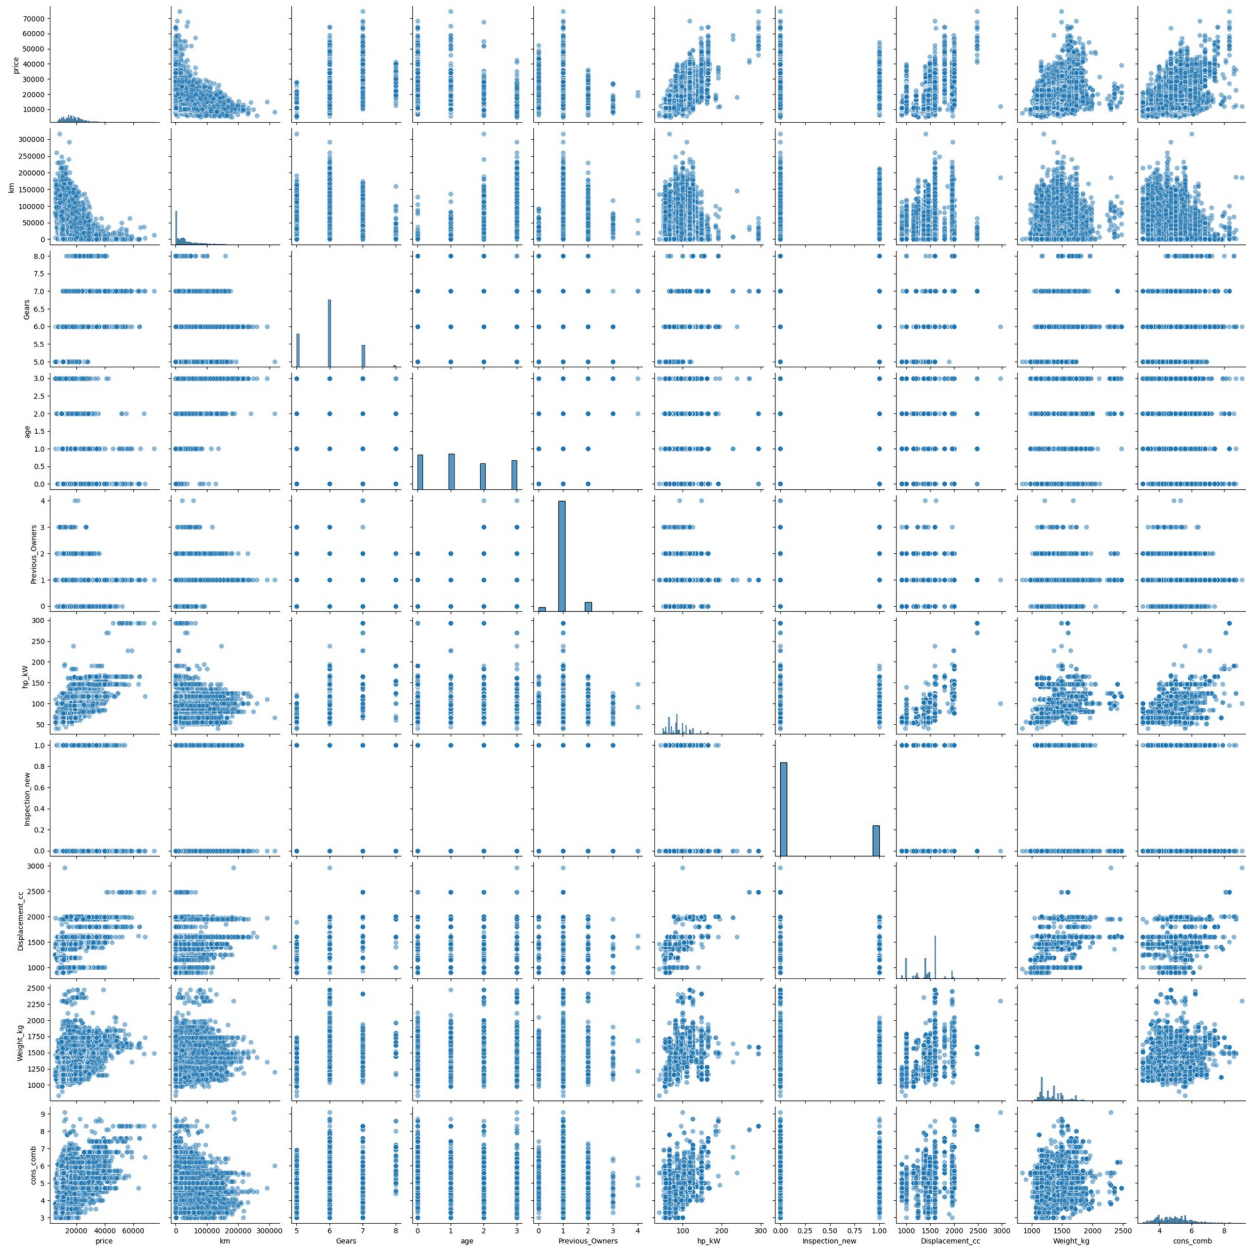
```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

<seaborn.axisgrid.PairGrid at 0x7caf248ad1e0>




we can observe that: 'km' will be good at polynomial regression and 'hp_kW', 'age', 'Weight_kg' have strong correlation will use at linear regression

```
# Calculate the correlation matrix
correlation_matrix = cleaned_df.corr()

# Create the heatmap
plt.figure(figsize=(8, 8)) # Set the size of the plot
sns.heatmap(correlation_matrix, annot=True, cmap='YlGnBu', fmt='.2f',
            linewidths=0.5)

# Show the plot
```

```
plt.title('Correlation Heatmap')
plt.show()
```




```
#best predictors for price in order
print(cleaned_df.corr().abs().nlargest(10,'price').index)

Index(['price', 'hp_kW', 'Gears', 'age', 'Weight_kg', 'km',
      'Displacement_cc',
      'cons_comb', 'Previous_Owners', 'Inspection_new'],
      dtype='object')
```



```
import matplotlib.pyplot as plt

plt.hist(cleaned_df['km'], bins=20, color='green', alpha=0.7)
plt.xlabel('km')
plt.ylabel('Frequency')
plt.title('Distribution of Horsepower')
plt.show()
```



linear regression

[illegible]

```
# Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

simple linear regression

```
y = cleaned_df['price']
X_simple = cleaned_df[['hp_kW']] # Best-correlated feature with price

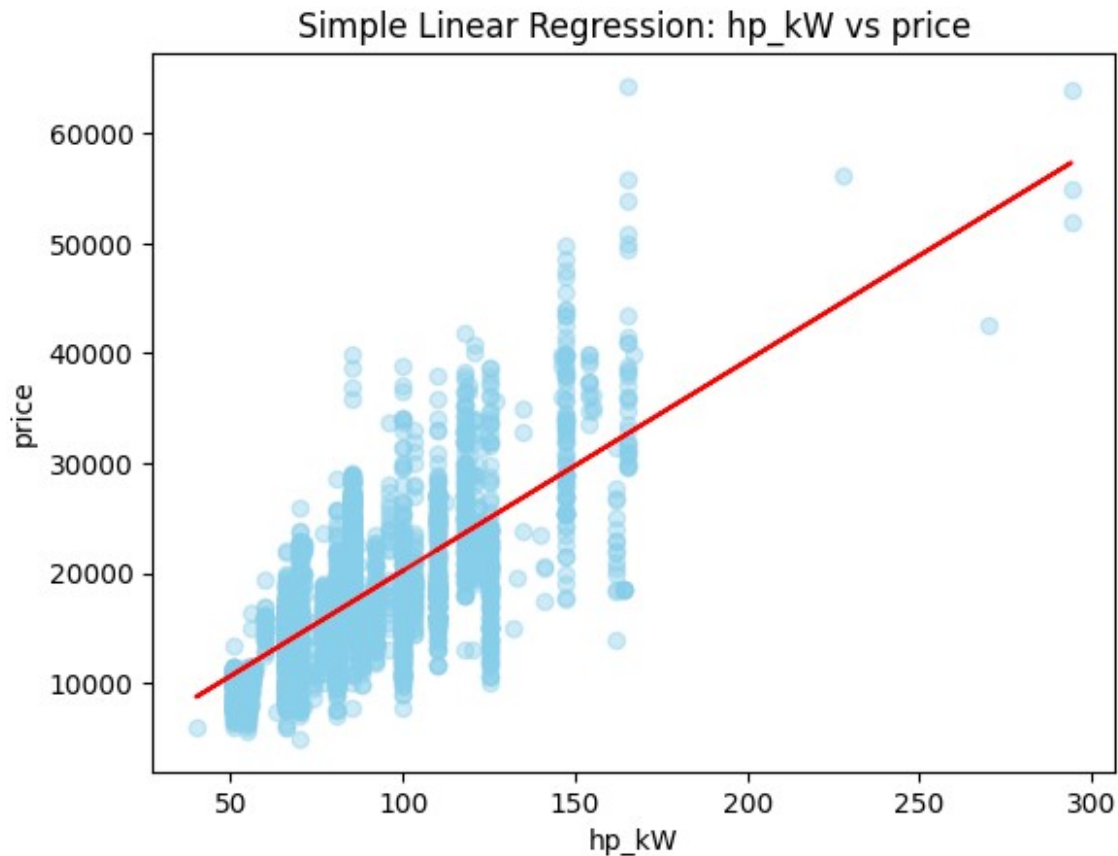
X_train_hp, X_test_hp, y_train_hp, y_test_hp =
train_test_split(X_simple, y, test_size=0.2, random_state=42)

model_simple = LinearRegression()
model_simple.fit(X_train_hp, y_train_hp)
y_pred_s = model_simple.predict(X_test_hp)

r2 = r2_score(y_test_hp, y_pred_s)
mse = mean_squared_error(y_test_hp, y_pred_s)
print("Simple Linear Regression:")
print("R^2 Score:", r2)
print("MSE:", mse)

# Visualization
plt.scatter(X_test_hp, y_test_hp, color='skyblue', alpha=0.4)
plt.plot(X_test_hp, y_pred_s, color='red')
plt.xlabel('hp_kW')
plt.ylabel('price')
plt.title('Simple Linear Regression: hp_kW vs price')
plt.show()

Simple Linear Regression:
R^2 Score: 0.5122324660161585
MSE: 26323633.657338142
```



```
# The coefficients in a dataframe
cdf =
pd.DataFrame(model_simple.coef_,X_simple.columns,columns=['Coef'])
cdf
```

```
      Coef
hp_kW 191.32341
```

multiple features linear regression

```
X_multi = cleaned_df[features]

X_train_m, X_test_m, y_train_m, y_test_m = train_test_split(X_multi,
y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train_m_scaled = scaler.fit_transform(X_train_m)
X_test_m_scaled = scaler.transform(X_test_m)

model_multi = LinearRegression()
model_multi.fit(X_train_m_scaled, y_train_m)
y_pred_m = model_multi.predict(X_test_m_scaled)
```

```

r2_m = r2_score(y_test_m, y_pred_m)
mse_m = mean_squared_error(y_test_m, y_pred_m)
print("Multiple Linear Regression:")
print("R^2 Score:", r2_m)
print("MSE:", mse_m)

```

```

Multiple Linear Regression:
R^2 Score: 0.757866741684256
MSE: 13067346.11076517

```

The coefficients in a dataframe

```

cdf = pd.DataFrame(model_multi.coef_, X_multi.columns, columns=['Coef'])
cdf

```

	Coef
km	-1535.268368
hp_kw	4391.514606
Displacement_cc	-685.228336
Weight_kg	912.094101
Previous_Owners	-13.227102
cons_comb	-938.892372
Gears	1684.476421
age	-2115.331738

#should be normal

```

residuals = y_test_m
sns.distplot(residuals, bins=30)

```

<ipython-input-218-3eb82e56d15b>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(residuals, bins=30)
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.

```

```


with pd.option_context('mode.use_inf_as_na', True):

```

```

<Axes: xlabel='price', ylabel='Density'>

```



```
comparison_df = pd.DataFrame({
    "Model": ["Simple Linear Regression", "Multiple Linear
Regression"],
    "R2 Score": [r2, r2_m],
    "MSE": [mse, mse_m]
})
```

```
comparison_df
```

	Model	R ² Score	MSE
0	Simple Linear Regression	0.512232	2.632363e+07
1	Multiple Linear Regression	0.757867	1.306735e+07

polynomial regression

```
y = cleaned_df['price']
X = cleaned_df[['km']]

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```

degrees = [2, 3, 4]
colors = ['green', 'blue', 'red']
mse_values = []

plt.figure(figsize=(10, 6))
plt.scatter(X_test, y_test, color='Silver', label='Actual Data',
alpha=0.6)

# Loop over each degree
for i, degree in enumerate(degrees):
    poly = PolynomialFeatures(degree=degree)
    X_train_poly = poly.fit_transform(X_train)
    X_test_poly = poly.transform(X_test)

    model = LinearRegression()
    model.fit(X_train_poly, y_train)
    y_pred = model.predict(X_test_poly)

    mse = mean_squared_error(y_test, y_pred)
    mse_values.append((degree, mse))

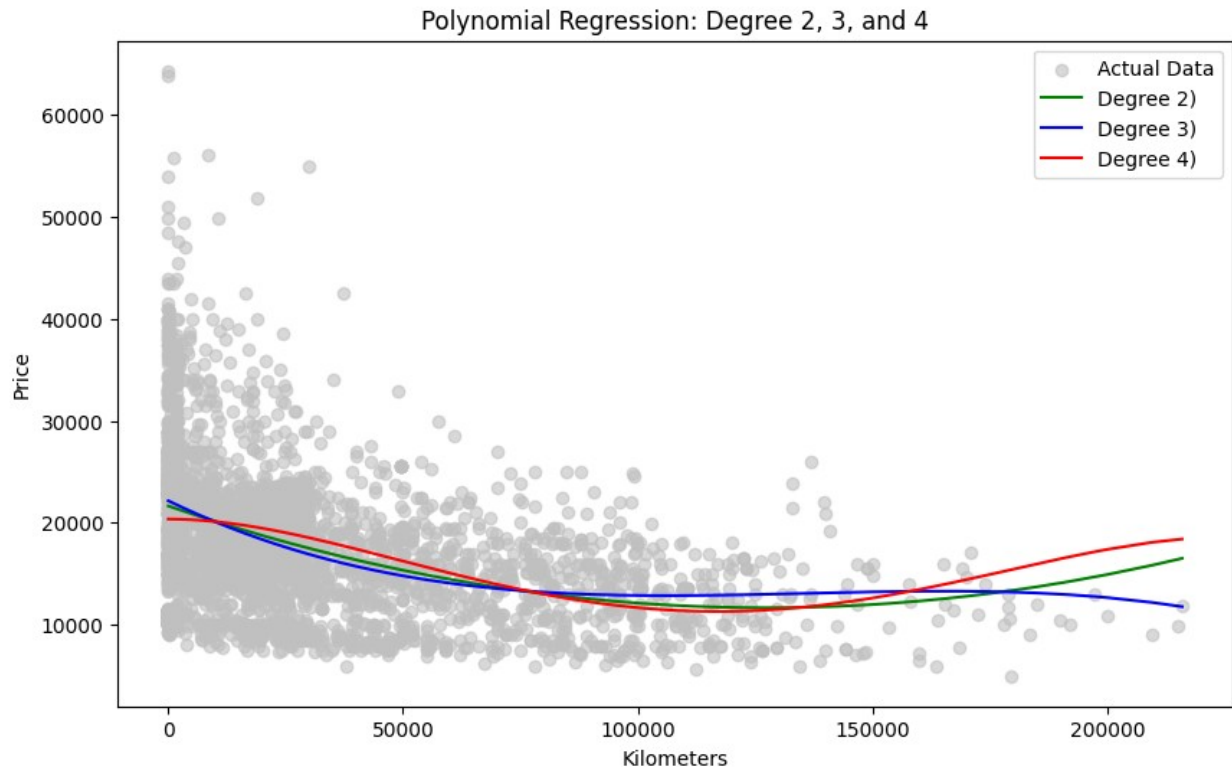
    # Sorting X for a smoother line plot
    sort_idx = X_test.values.flatten().argsort()
    X_sorted = X_test.values.flatten()[sort_idx]
    y_sorted = y_pred[sort_idx]

    plt.plot(X_sorted, y_sorted, color=colors[i], label=f'Degree
{degree}'))

plt.xlabel('Kilometers')
plt.ylabel('Price')
plt.title('Polynomial Regression: Degree 2, 3, and 4')
plt.legend()
plt.show()

# Print MSE values
for degree, mse in mse_values:
    print(f"Degree {degree}: MSE = {mse:.2f}")

```



Degree 2: MSE = 44508803.45
Degree 3: MSE = 44240412.40
Degree 4: MSE = 46067332.96

logistic regression

load the data

```
data =  
pd.read_csv("/kaggle/input/breast-cancer-wisconsin-data/data.csv")  
  
data.head()  
  
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/  
format.py:1458: RuntimeWarning: invalid value encountered in greater  
    has_large_values = (abs_vals > 1e6).any()  
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/format.py:14  
59: RuntimeWarning: invalid value encountered in less  
    has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals >  
0)).any()  
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/format.py:14  
59: RuntimeWarning: invalid value encountered in greater
```

```

    has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals >
0)).any()
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/format.py:14
58: RuntimeWarning: invalid value encountered in greater
    has_large_values = (abs_vals > 1e6).any()
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/format.py:14
59: RuntimeWarning: invalid value encountered in less
    has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals >
0)).any()
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/format.py:14
59: RuntimeWarning: invalid value encountered in greater
    has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals >
0)).any()

```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean
area_mean \					
0	842302	M	17.99	10.38	122.80
1001.0					
1	842517	M	20.57	17.77	132.90
1326.0					
2	84300903	M	19.69	21.25	130.00
1203.0					
3	84348301	M	11.42	20.38	77.58
386.1					
4	84358402	M	20.29	14.34	135.10
1297.0					

	smoothness_mean	compactness_mean	concavity_mean	concave
points_mean \				
0	0.11840	0.27760	0.3001	
0.14710				
1	0.08474	0.07864	0.0869	
0.07017				
2	0.10960	0.15990	0.1974	
0.12790				
3	0.14250	0.28390	0.2414	
0.10520				
4	0.10030	0.13280	0.1980	
0.10430				

	texture_worst	perimeter_worst	area_worst	
smoothness_worst \				
0 ...	17.33	184.60	2019.0	0.1622
1 ...	23.41	158.80	1956.0	0.1238
2 ...	25.53	152.50	1709.0	0.1444
3 ...	26.50	98.87	567.7	0.2098


```
4    ...           16.67           152.20           1575.0           0.1374
```

```
compactness_worst  concavity_worst  concave points_worst
symmetry_worst \
0           0.6656           0.7119           0.2654
0.4601
1           0.1866           0.2416           0.1860
0.2750
2           0.4245           0.4504           0.2430
0.3613
3           0.8663           0.6869           0.2575
0.6638
4           0.2050           0.4000           0.1625
0.2364
```

```
fractal_dimension_worst  Unnamed: 32
0           0.11890           NaN
1           0.08902           NaN
2           0.08758           NaN
3           0.17300           NaN
4           0.07678           NaN
```

```
[5 rows x 33 columns]
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 33 columns):
```


#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64

```
18 concavity_se          569 non-null    float64
19 concave points_se     569 non-null    float64
20 symmetry_se           569 non-null    float64
21 fractal_dimension_se  569 non-null    float64
22 radius_worst          569 non-null    float64
23 texture_worst         569 non-null    float64
24 perimeter_worst       569 non-null    float64
25 area_worst            569 non-null    float64
26 smoothness_worst      569 non-null    float64
27 compactness_worst     569 non-null    float64
28 concavity_worst       569 non-null    float64
29 concave points_worst  569 non-null    float64
30 symmetry_worst        569 non-null    float64
31 fractal_dimension_worst 569 non-null    float64
32 Unnamed: 32           0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

data cleaning

```
sns.heatmap(data.isnull())
```

```
<Axes: >
```



```
data.drop(["Unnamed: 32", "id"],axis = 1, inplace = True)
data.head()
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	M	17.99	10.38	122.80	1001.0	
1	M	20.57	17.77	132.90	1326.0	
2	M	19.69	21.25	130.00	1203.0	
3	M	11.42	20.38	77.58	386.1	
4	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave
points_mean \				
0	0.11840	0.27760	0.3001	
0.14710				
1	0.08474	0.07864	0.0869	
0.07017				
2	0.10960	0.15990	0.1974	
0.12790				
3	0.14250	0.28390	0.2414	

0.10520

4 0.10030 0.13280 0.1980

0.10430

	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	\
0	0.2419	...	25.38	17.33	184.60	
1	0.1812	...	24.99	23.41	158.80	
2	0.2069	...	23.57	25.53	152.50	
3	0.2597	...	14.91	26.50	98.87	
4	0.1809	...	22.54	16.67	152.20	

	area_worst	smoothness_worst	compactness_worst	concavity_worst	\
0	2019.0	0.1622	0.6656	0.7119	
1	1956.0	0.1238	0.1866	0.2416	
2	1709.0	0.1444	0.4245	0.4504	
3	567.7	0.2098	0.8663	0.6869	
4	1575.0	0.1374	0.2050	0.4000	

	concave	points_worst	symmetry_worst	fractal_dimension_worst
0		0.2654	0.4601	0.11890
1		0.1860	0.2750	0.08902
2		0.2430	0.3613	0.08758
3		0.2575	0.6638	0.17300
4		0.1625	0.2364	0.07678

[5 rows x 31 columns]

```
data.diagnosis = [1 if value == "M" else 0 for value in
data.diagnosis]
data.head()
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	1	17.99	10.38	122.80	1001.0	
1	1	20.57	17.77	132.90	1326.0	
2	1	19.69	21.25	130.00	1203.0	
3	1	11.42	20.38	77.58	386.1	
4	1	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave
points_mean \				
0	0.11840	0.27760	0.3001	
0.14710				
1	0.08474	0.07864	0.0869	
0.07017				
2	0.10960	0.15990	0.1974	
0.12790				
3	0.14250	0.28390	0.2414	
0.10520				
4	0.10030	0.13280	0.1980	
0.10430				

	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	\
0	0.2419	...	25.38	17.33	184.60	
1	0.1812	...	24.99	23.41	158.80	
2	0.2069	...	23.57	25.53	152.50	
3	0.2597	...	14.91	26.50	98.87	
4	0.1809	...	22.54	16.67	152.20	

	area_worst	smoothness_worst	compactness_worst	concavity_worst	\
0	2019.0	0.1622	0.6656	0.7119	
1	1956.0	0.1238	0.1866	0.2416	
2	1709.0	0.1444	0.4245	0.4504	
3	567.7	0.2098	0.8663	0.6869	
4	1575.0	0.1374	0.2050	0.4000	


	concave	points_worst	symmetry_worst	fractal_dimension_worst
0		0.2654	0.4601	0.11890
1		0.1860	0.2750	0.08902
2		0.2430	0.3613	0.08758
3		0.2575	0.6638	0.17300
4		0.1625	0.2364	0.07678

[5 rows x 31 columns]

data visualization

```
data["diagnosis"] = data['diagnosis'].astype("category", copy = False)
data["diagnosis"].value_counts().plot(kind = "bar")

<Axes: xlabel='diagnosis'>
```



```
y = data["diagnosis"]
X = data.drop(["diagnosis"], axis = 1)
y
0      1
1      1
2      1
3      1
4      1
..
564    1
565    1
566    1
567    1
568    0
Name: diagnosis, Length: 569, dtype: category
Categories (2, int64): [0, 1]
```

scaling

```
scaler = StandardScaler()
#fit
```

```
X_scaled = scaler.fit_transform(X)
X_scaled
array([[ 1.09706398, -2.07333501,  1.26993369, ...,  2.29607613,
         2.75062224,  1.93701461],
       [ 1.82982061, -0.35363241,  1.68595471, ...,  1.0870843 ,
        -0.24388967,  0.28118999],
       [ 1.57988811,  0.45618695,  1.56650313, ...,  1.95500035,
        1.152255   ,  0.20139121],
       ...,
       [ 0.70228425,  2.0455738 ,  0.67267578, ...,  0.41406869,
        -1.10454895, -0.31840916],
       [ 1.83834103,  2.33645719,  1.98252415, ...,  2.28998549,
        1.91908301,  2.21963528],
       [-1.80840125,  1.22179204, -1.81438851, ..., -1.74506282,
        -0.04813821, -0.75120669]])
```

split

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X_scaled, y,
test_size=0.30,random_state = 42)
```

train

```
from sklearn.linear_model import LogisticRegression
#Instantiate
lr = LogisticRegression()
#Train
lr.fit(X_train,y_train)
#Predict
y_pred = lr.predict(X_test)
y_pred
array([0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0,
0,
      1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0,
0,
      0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0,
0,
      1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1,
1,
      0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0,
0,
      1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1,
```

```
1,
    0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0,
0,
    0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0])
```

Confusion matrix

```
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test,y_pred)
print(f"Accuracy: {accuracy: 2f}")
```

Accuracy: 0.982456

```
from sklearn.metrics import classification_report
print(classification_report(y_pred,y_test))
```


	precision	recall	f1-score	support
0	0.98	0.99	0.99	107
1	0.98	0.97	0.98	64
accuracy			0.98	171
macro avg	0.98	0.98	0.98	171
weighted avg	0.98	0.98	0.98	171

```
from sklearn.metrics import confusion_matrix
confusion_matrix(y_test, y_pred)
```

```
array([[106,  2],
       [ 1, 62]])
```


```
from sklearn.metrics import ConfusionMatrixDisplay
```

```
disp =
ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test,
y_pred))
disp.plot()
plt.show()
```

```
import seaborn as sns

cm = confusion_matrix(y_test, y_pred)
# Plot the confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='g', cmap='Blues',
            xticklabels=['Predicted 0', 'Predicted 1'], yticklabels=['Actual 0',
            'Actual 1'])
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test,y_pred)
print(f"Accuracy: {accuracy:.2f}")
```

Accuracy: 0.98

```
from sklearn.metrics import classification_report
print(classification_report(y_pred,y_test))
```

	precision	recall	f1-score	support
0	0.98	0.99	0.99	107
1	0.98	0.97	0.98	64
accuracy			0.98	171
macro avg	0.98	0.98	0.98	171
weighted avg	0.98	0.98	0.98	171